Ethical AI Evaluation Report

Project: StressLens **Generated:** July 26, 2026

⚠️ **DATA QUALITY NOTICE:** The following data integrity issues were detected. Scores may be suppressed.

Cumulative Risk Volume

16.50

Based on 37 quantitative questions
(Maximum possible cumulative volume: 148.00)

Overall Average ERC (Mean): 0.27 / 4

Max Principle Average (Sensitivity): 0.44 / 4

ℹ️ *Overall risk promoted due to Max Principle Sensitivity override.*

(ERC values are normalized on a 0–4 scale)

Final Risk Label:
Minimal Risk

Ethical Tensions

2 IDENTIFIED

2 Accepted, 0 Under Review

Total Questions Assessed: 37 total (37 Quantitative, 0 Qualitative)

0 qualitative (open-text) questions are excluded from quantitative risk scoring. These questions provide narrative insights that complement the quantitative analysis.

Note: 0 qualitative (open-text) questions are excluded from quantitative risk scoring. Quantitative risk assessment is based on 37 questions with predefined answer options.

Ethical Principles Risk Overview

Ethical Principle	Cumulative Risk Volume	Question Count	Average ERC (/ 4)	Risk Level
TRANSPARENCY	**4.00**	9	0.44 / 4	Minimal Risk
HUMAN AGENCY & OVERSIGHT	**1.50**	8	0.19 / 4	Minimal Risk
TECHNICAL ROBUSTNESS & SAFETY	**2.00**	13	0.15 / 4	Minimal Risk
PRIVACY & DATA GOVERNANCE	**2.00**	8	0.25 / 4	Minimal Risk
DIVERSITY, NON-DISCRIMINATION & FAIRNESS	**0.00**	3	0.00 / 4	Minimal Risk
SOCIETAL & INTERPERSONAL WELL-BEING	**1.50**	6	0.25 / 4	Minimal Risk
ACCOUNTABILITY	**5.50**	14	0.39 / 4	Minimal Risk

Qualitative Analysis of Open-Text Responses

Methodology

In addition to quantitative scoring, expert evaluators provided open-text responses which have been analyzed.

The following analytical summaries synthesize expert observations with quantitative risk assessments.

Transparency

[MOCK DATA] The transparency measures are present but need better communication strategies to users.

Disclaimer: The qualitative insights presented above are intended to provide contextual understanding.

Executive Summary

- The assessment identified a Cumulative Risk Volume of 43 across 10 evaluated quantitative questions.
- The Normalized Average ERC is 2.5 / 4, classified as Elevated Risk.
- Data privacy and unmitigated bias are the primary drivers of this risk level.
- The system shows strong documentation practices in technical robustness.
- Overall, significant interventions are required before production deployment.

Ethical Tensions

Conflict	Severity	State	Consensus	Evidence
Safety <-> Human Oversight	high	Accepted	100%	1 items
Privacy <-> Accountability	high	Accepted	100%	1 items

Methodology & Evidence Coverage

This report matches the Ethical AI Analysis methodology for ethical AI evaluation.

Risk Calculation

- **Question Risk (ERC):** Importance (0-4) x Unmitigated Ethical Risk (0-1).
- **Cumulative Risk Volume:** Sum of all ERC contributions. Used to understand total magnitude of risk for the project.
- **Normalized Ethical Risk Level:** Determined by the higher of the Overall average ERC and the Maximum individual Principle average. This promotes the risk label if any single principle (e.g. Accountability) identifies a critical failure, even if the overall project volume is low.

Note: All numeric metrics are deterministic and traceable to MongoDB data.

Appendix: Evaluators

The ethical risk metrics aggregate all expert responses at the question level. While individual experts may differ, the normalized score reflects the combined judgment across evaluators rather than any single opinion.

Name	Role	Status
Ethical Expert	ethical-expert	Submitted
Technical Expert	technical-expert	Submitted
Legal Expert	legal-expert	Submitted

Improvement Recommendations

Short-Term (0-3 Months)

- [MOCK DATA] Implement explicit consent
- [MOCK DATA] Update privacy policy

Medium-Term (3-12 Months)

- [MOCK DATA] Conduct algorithmic bias audit
- [MOCK DATA] Establish human-in-the-loop protocols

Long-Term (12+ Months)

- [MOCK DATA] Create continuous monitoring framework
- [MOCK DATA] Form external ethics board

Conclusion

Overall, the system demonstrates foundational awareness of ethical AI principles but lacks robust implementation.

The most critical weakness is the handling of sensitive biometric data, while the strength lies in technical safety.

Future governance should focus on strict access controls and regular third-party audits.